

CS6700: Tutorial 3

Value & Policy Iteration

Gautham Govind A, EE19B022


```
import numpy as np
from enum import Enum
import copy
```

Consider a standard grid world, where only 4 (up, down, left, right) actions are allowed and the agent deterministically moves accordingly, represented as below. Here yellow is the start state and white is the goal state.

Say, we define our MDP as:

- S: 121 (11 x 11) cells
- A: 4 actions (up, down, left, right)
- P: Deterministic transition probability
- R: -1 at every step
- gamma: 0.9

Our goal is to find an optimal policy (shown in right).



```
# Above grid is defined as below:
# - 0 denotes an navigable tile
# - 1 denotes an obstruction/wall
# - 2 denotes the start state
# - 3 denotes an goal state
```

```
# Note: Here the upper left corner is defined as (0, 0)
```

```
#         and lower right corner as (m-1, n-1)

# Optimal Path: RIGHT RIGHT UP UP LEFT LEFT UP UP UP UP UP UP LEFT
LEFT DOWN DOWN LEFT LEFT
```

```
GRID_WORLD = np.array([
    [1, 1, 1, 1, 1, 1, 1, 1, 1, 1],
    [1, 0, 0, 0, 0, 0, 0, 0, 1, 1],
    [1, 1, 1, 0, 1, 0, 1, 1, 1, 1],
    [1, 3, 0, 0, 1, 0, 1, 0, 1, 1],
    [1, 1, 1, 1, 1, 0, 1, 0, 1, 1],
    [1, 0, 0, 0, 1, 0, 1, 0, 0, 1],
    [1, 1, 1, 0, 1, 0, 1, 0, 1, 1],
    [1, 0, 0, 0, 0, 0, 0, 0, 1, 1],
    [1, 1, 1, 1, 1, 1, 1, 0, 1, 1],
    [1, 0, 0, 0, 0, 2, 0, 0, 0, 1],
    [1, 1, 1, 1, 1, 1, 1, 1, 1, 1]
])
```

Actions

```
class Actions(Enum):
    UP = (0, (-1, 0)) # index = 0, (xaxis_move = -1 and yaxis_move =
0)
    DOWN = (1, (1, 0)) # index = 1, (xaxis_move = 1 and yaxis_move =
0)
    LEFT = (2, (0, -1)) # index = 2, (xaxis_move = 0 and yaxis_move =
-1)
    RIGHT = (3, (0, 1)) # index = 3, (xaxis_move = 0 and yaxis_move =
-1)

    def get_action_dir(self):
        _, direction = self.value
        return direction

    @property
    def index(self):
        indx, _ = self.value
        return indx

    @classmethod
    def from_index(cls, index):
        action_index_map = {a.index: a for a in cls}
        return action_index_map[index]

# How to use Action enum
for a in Actions:
    print(f"name: {a.name}, action_id: {a.index}, direction_to_move:
{a.get_action_dir()}")
```

```
print("\n-----\n")
```

```
# find action enum from index 0
```

```
a = Actions.from_index(0)
```

```
print(f"0 index action is: {a.name}")
```

```
name: UP, action_id: 0, direction_to_move: (-1, 0)
```

```
name: DOWN, action_id: 1, direction_to_move: (1, 0)
```

```
name: LEFT, action_id: 2, direction_to_move: (0, -1)
```

```
name: RIGHT, action_id: 3, direction_to_move: (0, 1)
```

```
-----
```

```
0 index action is: UP
```

Policy

```
class BasePolicy:
```

```
    def update(self, *args):
```

```
        pass
```

```
    def select_action(self, state_id: int) -> int:
```

```
        raise NotImplemented
```

```
class DeterministicPolicy(BasePolicy):
```

```
    def __init__(self, actions: np.ndarray):
```

```
        # actions: its a 1d array (|S| size) which contains action for  
each state
```

```
        self.actions = actions
```

```
    def update(self, state_id, action_id):
```

```
        assert state_id < len(self.actions), f"Invalid state_id  
{state_id}"
```

```
        assert action_id < len(Actions), f"Invalid action_id {action_id}"
```

```
        self.actions[state_id] = action_id
```

```
    def select_action(self, state_id: int) -> int:
```

```
        assert state_id < len(self.actions), f"Invalid state_id  
{state_id}"
```

```
        return self.actions[state_id]
```

Environment

```
class Environment:
```

```
    def __init__(self, grid):
```

```
        self.grid = grid
```

```
        m, n = grid.shape
```

```
        self.num_states = m*n
```

```
    def xy_to_posid(self, x: int, y: int):
```

```
        _, n = self.grid.shape
```

```

    return x*n + y

def posid_to_xy(self, posid: int):
    _, n = self.grid.shape
    return (posid // n, posid % n)

def isvalid_move(self, x: int, y: int):
    m, n = self.grid.shape
    return (x >= 0) and (y >= 0) and (x < m) and (y < n) and
(self.grid[x, y] != 1)

def find_start_xy(self) -> int:
    m, n = self.grid.shape
    for x in range(m):
        for y in range(n):
            if self.grid[x, y] == 2:
                return (x, y)
    raise Exception("Start position not found.")

def find_path(self, policy: BasePolicy) -> str:
    max_steps = 50
    steps = 0

    P, R = self.get_transition_prob_and_expected_reward()
    num_actions, num_states = R.shape
    all_possible_state_posids = np.arange(num_states)

    path = ""
    curr_x, curr_y = self.find_start_xy()
    while (self.grid[curr_x, curr_y] != 3) and (steps < max_steps):
        curr_posid = self.xy_to_posid(curr_x, curr_y)
        action_id = policy.select_action(curr_posid)
        next_posid = np.random.choice(
            all_possible_state_posids, p=P[action_id, curr_posid])
        action = Actions.from_index(action_id)
        path += f" {action.name}"
        curr_x, curr_y = self.posid_to_xy(next_posid)
        steps += 1
    return path

def get_transition_prob_and_expected_reward(self): #  $P(s_{next} | s, a)$ ,  $R(s, a)$ 
    m, n = self.grid.shape
    num_states = m*n
    num_actions = len(Actions)
    P = np.zeros((num_actions, num_states, num_states))
    R = np.zeros((num_actions, num_states))
    for a in Actions:
        for x in range(m):

```

```

    for y in range(n):
        xmove_dir, ymove_dir = a.get_action_dir()
        xnew, ynew = x + xmove_dir, y + ymove_dir # find the new
co-ordinate after the action a

        posid = self.xy_to_posid(x, y)
        new_posid = self.xy_to_posid(xnew, ynew)

        if self.grid[x, y] == 3:
            # the current state is a goal state
            P[a.index, posid, posid] = 1
            R[a.index, posid] = 0
        elif (self.grid[x, y] == 1) or (not self.isvalid_move(xnew,
ynew)):
            # the current state is a block state or the next state is
invalid
            P[a.index, posid, posid] = 1
            R[a.index, posid] = -1
        else:
            # action a is valid and goes to a new position
            P[a.index, posid, new_posid] = 1
            R[a.index, posid] = -1
    return P, R

```

Policy Iteration

```

def policy_evaluation(P: np.ndarray, R: np.ndarray, gamma: float,
                    policy: BasePolicy, theta: float,
                    init_V: np.ndarray=None):
    num_actions, num_states = R.shape

    # Please try different starting point for V you will find it will
always
# converge to the same V_pi value.
    if init_V is None:
        init_V = np.zeros(num_states)
    V = copy.deepcopy(init_V)

    delta = 100.0
    while delta > theta:
        delta = 0.0
        for state_id in range(num_states):
            action_id = policy.select_action(state_id)
            v_old = V[state_id]
            # Following equation is a different way of writing the same
equation given in the slide.
            # Note here R is an expected reward term.
            V[state_id] = R[action_id, state_id] + gamma *
np.dot(P[action_id, state_id], V)
            delta = max(delta, abs(V[state_id] - v_old))

```

```
return V
```

```
def policy_improvement(P: np.ndarray, R: np.ndarray, gamma: float,
                      policy: BasePolicy, V: np.ndarray):
    num_actions, num_states = R.shape
    policy_stable = True
    for state_id in range(num_states):
        old_action_id = policy.select_action(state_id)

        # your code here
        new_action_id = np.argmax(R[:, state_id] + gamma * np.dot(P[:,
state_id], V)) # update new_action_id based on the value function.

        policy.update(state_id, new_action_id)
        if old_action_id != new_action_id:
            policy_stable = False
    return policy_stable
```

```
def policy_iteration(P: np.ndarray, R: np.ndarray, gamma: float,
                    theta: float=1e-3, init_policy: BasePolicy =
None):
    num_actions, num_states = R.shape

    # Please try exploring different policies you will find it will
    always
    # converge to the same optimal policy for valid states.
    if init_policy is None:
        # Say initial policy = all up actions.
        init_policy = DeterministicPolicy(actions=np.zeros(num_states,
dtype=int))

    # creating a copy of a initial policy
    policy = copy.deepcopy(init_policy)
    policy_stable = False
    while not policy_stable:
        V = policy_evaluation(P, R, gamma, policy, theta)
        policy_stable = policy_improvement(P, R, gamma, policy, V)
    return policy, V
```

Value Iteration

```
# Directly find the optimal value function
```

```
def get_optimal_value(P: np.ndarray, R: np.ndarray, gamma: float,
                     theta: float, init_V: np.ndarray=None):
    num_actions, num_states = R.shape

    # Please try different starting point for V you will find it will
    always
```

```

# converge to the same V_star value.
if init_V is None:
    init_V = np.zeros(num_states)
V = copy.deepcopy(init_V)

delta = 100.0
while delta > theta:
    delta = 0.0
    for state_id in range(num_states):
        v_old = V[state_id]
        q_sa = np.zeros(num_actions)
        for a in Actions:
            q_sa[a.index] = R[a.index, state_id] + gamma *
np.dot(P[a.index, state_id], V)
        V[state_id] = np.max(q_sa)
        delta = max(delta, abs(V[state_id] - v_old))
    return V

def value_iteration(P: np.ndarray, R: np.ndarray, gamma: float,
                    theta: float=1e-3, init_V: np.ndarray=None):
    V_star = get_optimal_value(P, R, gamma, theta, init_V)

    num_actions, num_states = R.shape
    policy = DeterministicPolicy(actions=np.zeros(num_states,
dtype=int))
    for state_id in range(num_states):
        # Your code here
        action_id = np.argmax(R[:, state_id] + gamma * np.dot(P[:,
state_id], V_star)) # update the action_id based on V_star

        policy.update(state_id, action_id)

    return policy, V_star

```

Experiments

```

def is_same_optimal_value(V1, V2, diff_theta=1e-3):
    diff = np.abs(V1 - V2)
    return np.all(diff < diff_theta)

```

```

seed = 0
np.random.seed(seed)

```

```

gamma = 0.9
theta = 1e-5

```

```

env = Environment(GRID_WORLD)
P, R = env.get_transition_prob_and_expected_reward()

```

Exp 1: Using Policy iteration algorithm find the optimal path from start to goal position

```
# # Start with random choice of init_policy.  
# One such choice could be: init_policy = np.ones(env.num_states,  
dtype=int)  
# Start with your own choice of init_policy  
init_policy = DeterministicPolicy(actions=2*np.ones(env.num_states,  
dtype=int))
```

```
pitr_policy, pitr_V_star = policy_iteration(P, R, gamma, theta=theta,  
init_policy=init_policy)  
pittr_path = env.find_path(pitr_policy)  
print(pitr_path)
```

```
RIGHT RIGHT UP UP LEFT LEFT UP UP UP UP UP UP LEFT LEFT DOWN DOWN  
LEFT LEFT
```

Exp 2: Using value iteration algorithm find the optimal path from start to goal position

```
vittr_policy, vittr_V_star = value_iteration(P, R, gamma, theta=theta)  
vittr_path = env.find_path(vitr_policy)  
print(vitr_path)
```

```
RIGHT RIGHT UP UP LEFT LEFT UP UP UP UP UP UP LEFT LEFT DOWN DOWN  
LEFT LEFT
```

Exp 3: Compare the optimal value function of policy iteration and value iteration algorithm

```
is_same_optimal_value(pitr_V_star, vittr_V_star)
```

```
True
```

Exp 4: Using initial guess for V as random values, find the optimal value function using policy evaluation and compare it with the optimal value function

```
# Start with random choice of init_V.  
# One such choice could be: init_V = np.random.randn(env.num_states)  
# Another choice could be: init_V = 10*np.ones(env.num_states)  
# Start with your own choice of init_V  
init_V = 5*np.ones(env.num_states) # your choice
```

```
V_star = policy_evaluation(P, R, gamma, pitr_policy, theta, init_V)  
is_same_optimal_value(pitr_V_star, V_star)
```

```
True
```

Exp 5: Using initial guess for V as random values, find the optimal value function using get_optimal_value and compare it with the optimal value function

```
# Start with random choice.  
# One such choice could be: init_V = np.random.randn(env.num_states)  
# Another choice could be: init_V = 10*np.ones(env.num_states)  
# Start with your own choice of init_V  
init_V = 3*np.ones(env.num_states)
```



```
V_star = get_optimal_value(P, R, gamma, theta, init_V)
is_same_optimal_value(vitr_V_star, V_star)
```

True

Exp Optional: Try changing the grid by adding multiple paths to the goal state and check if our policy_iteration or value_iteration algorithm is able to find optimal path. Redo the above experiments.

```
• 1 way to add another path would be GRID_WORLD[4, 1] = 0
GRID_WORLD[4, 1] = 0
env = Environment(GRID_WORLD)
P, R = env.get_transition_prob_and_expected_reward()

init_policy = DeterministicPolicy(actions=2*np.ones(env.num_states,
dtype=int))

pitr_policy, pitr_V_star = policy_iteration(P, R, gamma, theta=theta,
init_policy=init_policy)
pitr_path = env.find_path(pitr_policy)
print(pitr_path)

RIGHT RIGHT UP UP LEFT LEFT LEFT LEFT UP UP LEFT LEFT UP UP

vitr_policy, vitr_V_star = value_iteration(P, R, gamma, theta=theta)
vitr_path = env.find_path(vitr_policy)
print(vitr_path)

RIGHT RIGHT UP UP LEFT LEFT LEFT LEFT UP UP LEFT LEFT UP UP

is_same_optimal_value(pitr_V_star, vitr_V_star)

True

init_V = 7*np.ones(env.num_states) # your choice

V_star = policy_evaluation(P, R, gamma, pitr_policy, theta, init_V)
is_same_optimal_value(pitr_V_star, V_star)

True

init_V = 9*np.ones(env.num_states)

V_star = get_optimal_value(P, R, gamma, theta, init_V)
is_same_optimal_value(vitr_V_star, V_star)

True
```